

Sec. 5 Experiments

Table 1: Headline cells, tail-500 over 5 seeds, decaying step-size. *full* sees every group every step. The fixed multiplier m/K diverges whenever observation rates are non-uniform.

Instance	GAVOT	unif. avg	m/K	full
Adult, prevalence ($\nu=.60$)	0.2269	0.2479	0.670	0.2073
Adult, aligned ($\nu=0$)	0.2075	0.2077	div.	0.2073
IMDb, skewed ($\nu=.31$)	86.06	91.63	8504	82.95
IMDb, aligned ($\nu=0$)	84.28	86.37	div.	82.95

1 EXPERIMENTS

Instances. *Adult* income classification [33]: logistic regression ($d=109$ with intercept, cross-entropy, $\eta=0.1$) on 100 race-homogeneous groups of 30 samples, the number of groups per race proportional to its prevalence (White/Black/Asian-Pac./Amer-Ind./Other = 85/10/3/1/1); the importance target is the fairness target itself, $1/5$ per race divided equally among its groups, so a group from the smallest race carries $p_i = 0.2$. *IMDb-Wiki* age regression [34]: linear regression of age on 128-d ResNet face embeddings ($\eta=10^{-2}$) over the first 100 identities with 30 images, with the cubic importance target $p_i \propto (101 - i)^3$. Observation rates are swept through a one-parameter family: on Adult $r \propto p^\beta$, $\beta \in \{0, \frac{1}{4}, \frac{1}{2}, \frac{3}{4}, 1\}$, from uniform observation (prevalence regime, $\nu=.60$: the five groups of the three smallest races hold 60% of the target against $r_i \approx K/m$) to importance-aligned ($\nu=0$); on IMDb-Wiki $r_i \propto i^\beta$, $\beta \in \{0, \frac{1}{2}, 1, \frac{3}{2}, 2, 3\}$, from uniform observation ($\nu=.31$, nine groups) to the mirror image of p ($\nu=.88$, 41 groups), plus the aligned $r_i \propto 101 - i$ ($\nu=0$). Eleven instances, ν from 0 to .88.

Protocol. Each step observes $K=3$ groups drawn without replacement with probabilities $\propto r$, which induces q over the $\binom{100}{3}$ compositions; q is estimated from 10^6 draws of the sampler and the plan from (7) on the membership mask, once per instance. Each observed group takes $H=5$ local steps on its own data before the weighted average, the standard practical deviation from the one-step analysis, applied to every rule alike since all rules consume the same local updates from the same draw A^t . Rules: *GAVOT*; *group-blind averaging*, the uniform average over the observed groups (Proposition 2 with $\tilde{p}_i = r_i/K$); the *fixed multiplier* m/K , the rescaling $(m/K)p_i$ that presumes uniform observation; and *full*, every group every step weighted by p , the reference $\Phi(p)$. Stepsize $\eta_t = \eta/(1 + t/1000)$ unless stated, 4000 steps, 5 seeds; we report F_p averaged over the last 500 steps, mean $\pm$ std over seeds. On both instances $\arg \min F_w$ is available to machine precision (closed form for least squares, Newton for the logistic loss), so the floors $\Phi(p), \Phi(\tilde{p}), \Phi(\bar{p})$ of Theorem 8 are exact rather than estimated, and in every instance with $\nu = 0$ the iteration (7) converges to a row error below $3 \cdot 10^{-8}$, certifying the remaining inequalities of (5).

The least represented group. Fig. 2 reads the same runs on the group the target exists for. On Adult, trans-




Figure 1: IMDb-Wiki severity sweep. Solid, measured tail-500 (± 1 std); dashed, the exact floor $\Phi(\cdot)$ each rule converges to. Left, constant stepsize: the optimization residual swamps the alignment gain and transport losses everywhere, even where the floors are seven units apart. Right, decaying stepsize: the residual collapses and the measured curves track their floors, so transport wins until the floors close at high ν . The sign is set by which of the two terms of (9) dominates, not by the regime label.

port lowers the loss on race Other at every skew and under both stepsizes: .073 against .119 for group-blind averaging on the prevalence instance (full coverage .031), and .031 against .035 on the aligned one, where it matches full coverage. Amer-Indian, the other single-group race, behaves the same way. On IMDb-Wiki the top-importance identities show the crossover of Table 2 one step earlier: under decay transport wins at $\nu \leq .70$ (77.0 against 86.4 at $\nu=.31$, full 72.2; 73.6 against 79.9 when aligned) and loses at the three most infeasible skews, while at constant stepsize it wins only at $\nu=.31$. The overall objective and the rare-group loss therefore tell the same story, with the rare group the more sensitive of the two.

E1: the headline cells. Table 1 reports the four instances a practitioner would recognize. Transport wins on three: Adult under prevalence-proportional availability, 0.2269 against 0.2479 ($t=45$), and both IMDb cells, 86.06 against 91.63 ($t=21$) and 84.28 against 86.37 ($t=13$). On Adult with importance-aligned availability the two coincide at the oracle value, which is what Proposition 2 predicts when $\tilde{p} = p$: there is nothing to correct and transport does not manufacture a difference. The fixed multiplier diverges on both aligned instances and is far off on the other two, so naive rescaling is no alternative to solving (4).

E2: the crossover, and where the sign comes from. Table 2 and Fig. 1 are the paper’s main evidence. On Adult, whose five skews are omitted from the table, the measured gain tracks Δ_{floor} within 4% at three of five (.0134 against .0129, .0050 against .0049, .0002 against .0003) and exceeds it at the two most infeasible (.0210 against .0110, .0181 against .0144), group-blind averaging sitting further above its floor. On IMDb-Wiki the gain is uniformly below the floor by an excess optimization residual that is close to constant in β , mean 1.9 in MSE units, and the measured sign flips precisely where Δ_{floor} crosses that value, between $\beta=1.5$ (2.56, gain




Figure 2: Loss on the least represented group, the metric the fairness target is about. Top, Adult: cross-entropy on race Other, one group of 30 out of 100, to which the category-uniform target assigns $p_i = 0.2$; the axis is the observation-rate exponent, $\beta=1$ aligned, $\beta=0$ uniform (prevalence). Bottom, IMDb-Wiki: MSE on the 20 highest-importance identities, the least observed ones once $\beta > 0$, against ν . Left panels constant stepsize, right panels decaying. Same protocol as the tables ($K=3$, $H=5$, 4000 steps, tail-500 mean $\pm$ std over 5 seeds); lower is better. Curves are the group-wise entries of the same runs as Table 1 and Fig. 1.

+0.58) and $\beta=2.0$ (1.62, gain -0.03). This is Theorem 8 read at finite T : the sign of the comparison is a quantity one computes in advance.

E3: the ℓ_1 bound is too loose to be the criterion. The distances $\|p - \hat{p}\|_1$ and $\|p - \tilde{p}\|_1$ are also computed on every instance, and $\|p - \hat{p}\|_1 < \|p - \tilde{p}\|_1$ holds in all eleven, including the three where transport loses or ties. The ℓ_1 form of Theorem 8 therefore never predicts a loss and cannot serve as the criterion; the exact floors can.

E4: the residual is a stepsize artifact. The left panel of Fig. 1 runs the identical instances at constant stepsize. There transport loses on all six IMDb skews, by as much as 7.9, even at $\nu=.31$ where the floors are 6.99 apart. The transport weights $Y[:, j]$ concentrate mass on rarely observed groups, so although Lemma 6 bounds the second moment identically for both rules, the variance about the mean is larger, and at constant stepsize the iterates sit on a noise floor proportional to it. Under $\eta_t = \eta/(1 + t/1000)$ that floor decays, the measured curves fall onto their dashed floors, and every sign reversal attributable to the stepsize disappears. The residual is a property of the schedule, not of the estimator.

E5: the decomposition in closed form. The mirrored IMDb-Wiki instance ($\nu=.88$) is where Theorem 8 can be read off exactly. The Sinkhorn limit truncates the target at the observation-rate ceiling, $\hat{p}_i \approx \min(p_i, r_i)$ up to renormalization (maximum deviation .023), with $\|p - \hat{p}\|_1 = 1.58$ of a possible 2. The frontier is least squares, so $\Phi(p) = 82.9$, matching the measured full plateau 82.95, and $\Phi(\hat{p}) = 105.9$:

Table 2: IMDb-Wiki severity sweep, decaying stepsize. Δ_{floor} is the alignment gain of Theorem 8, *gain* the measured advantage of transport, t its Welch statistic (population std). The excess residual transport pays averages 1.9, and the sign flips exactly where Δ_{floor} crosses it. The five Adult skews all sit above their crossover and are omitted.

β	ν	Δ_{floor}	gain	t
0	.31	6.99	+5.57	20.7
0.5	.59	5.71	+3.46	11.3
1.0	.70	3.67	+2.05	5.4
1.5	.77	2.56	+0.58	1.1
2.0	.82	1.62	-0.03	-0.1
3.0	.88	0.02	-2.28	-7.9

the alignment floor is 23.0 and the realized bias equals its exact expression $|(p - \hat{p})^\top L(\theta^\dagger)| = 30.7$, where L collects the per-group risks, while the ℓ_1 form $2M\|p - \hat{p}\|_1$ is conservative since $M \geq \max_i L_i(\theta^\dagger) \approx 454$. Group-blind averaging has the same floor to two decimals, $\Phi(\tilde{p}) = 105.9$: both surrogates starve the same groups, so Δ_{floor} vanishes and the entire measured gap between the two rules at this endpoint of Table 2 is residual, which is why it is the one instance where transport still loses by a significant margin under decay ($\beta=2$, where the floors are 1.6 apart, is a tie).

E6: risk aversion, and where it does not help. Sweeping (α, γ) on a 10×10 grid, the overall objective is minimized on the $\gamma = 1$ edge on every instance, which is $\kappa(1, \gamma) = 1$ made visible and confirms that risk aversion costs the average. The one place the CVaR layer earns its keep is the worst race on Adult under prevalence availability, where GAVOT+CVaR at $(0.2, 0.7)$ reaches 0.3624 against 0.3679 for GAVOT and 0.3994 for group-blind averaging. Per race, transport improves four of five (Black .179 vs .200, Asian-Pac .368 vs .399, Am-Ind .175 vs .189, Other .073 vs .119) and regresses only the majority group (.339 vs .332), which is what reweighting toward a category-uniform target is supposed to do.