

# Group-Fair Stochastic Gradient Descent via Masked Optimal Transport

Herlock (SeyedAbolfazl) Rahimi     Dionysis Kalogerias<sup>1</sup>

Department of Electrical and Computer Engineering, Yale University, New Haven, CT, USA  
{herlock.rahimi, dionysis.kalogerias}@yale.edu

## ABSTRACT

*Minibatch SGD is group-blind. Which groups appear in a batch is decided by how often they occur in the data, not by how much they are meant to count in the objective, and when the two disagree the iterates converge to a prevalence-weighted surrogate rather than to the target risk. The gap is a bias, so no amount of training removes it. A minibatch is a partial-participation event over groups, which makes group-fair training an instance of federated aggregation under availability skew, and we carry the masked optimal transport construction of FedAVOT across. Aligning the batch-composition distribution  $q$  with the target importance distribution  $p$  produces per-batch group weights whose induced stochastic sub-gradient is exactly unbiased for the target risk, giving the standard  $O(1/\sqrt{T})$  rate in the nonsmooth convex regime with a constant that does not depend on how many groups a batch happens to cover. Exact alignment is feasible precisely under a Hall-type condition, which reduces to the per-group requirement that importance not exceed observation rate; past it Sinkhorn returns an information projection and a residual we compute in closed form rather than bound. That closed form turns the excess risk into two terms, an alignment gain fixed by the instance and an optimization residual paid by the estimator, and a severity sweep over eleven availability skews on Adult and IMDb-Wiki shows the two crossing: transport wins while the gain exceeds the residual and loses past it, at a threshold the decomposition locates in advance. The residual is a stepsize artifact, not a property of the method, and a decaying schedule removes it on every instance where it had reversed the sign. A mean-CVaR layer composed on top reaches the worst group at rate  $O(\kappa(\alpha, \gamma)/\sqrt{T})$ .*

**Index Terms**— Group fairness, optimal transport, stochastic gradient descent, distribution shift, conditional value at risk.

## 1 INTRODUCTION

A classifier is rarely asked to minimize one risk. It is asked to minimize one while behaving acceptably on sub-populations the training distribution under-represents: a diagnostic model across patient groups, a recognition system across demographic strata [1, 2, 3]. The standard response is to change the objective, replacing the empirical risk by  $F_p(\theta) = \sum_i p_i f_i(\theta)$  with  $p$  the weighting one actually wants, uniform, inverse to prevalence, or externally dictated [4, 5].

This paper is about a gap between writing that objective down and optimizing it. Minibatch SGD does not see  $p$ : it sees a batch, whose group composition is governed by prevalence and by the sampler, so averaging the per-example gradients descends a different function. The estimator is not

merely noisy, it is biased, by a fixed vector that survives  $T \rightarrow \infty$  and any stepsize schedule. Rebalancing the sampler is the usual and only partial fix: a group whose prevalence is below the resolution of the batch size cannot be sampled into every batch however the sampler is tilted, and upscaling the few examples that do appear reintroduces an amplitude distortion that destabilizes the iterates [6].

**Groups as clients.** The same mismatch is the central object of a different literature. In federated learning the *availability distribution*  $q$  over which clients participate in a round rarely matches the *importance distribution*  $p$  that defines the objective, and FedAvg consequently optimizes a surrogate whose weights are induced by the availability process rather than chosen [7, 8, 9]. FedAVOT [6] removes that mismatch by posing aggregation as a masked optimal transport (MOT) problem: find a coupling whose column marginal is the availability law and whose row marginal is the target  $p$ , supported only on client-round pairs that actually occurred. The observation driving this paper is that a minibatch is a participation event: the groups represented in  $B^t$  are the active set  $S^t$ , its law is  $q$ , and the weighting one wants is  $p$ . Group-fair SGD is federated aggregation with the federation collapsed onto one machine. We call the method GAVOT and keep the FedAVOT label in figures, the aggregation rule being unchanged.

### Contributions.

1. *The reduction* (Sec. 2). Group-blind minibatch SGD reduces, under uniform-within-batch averaging, to agnostic FedAvg with one local step, so the surrogate identity of [7] names the function SGD actually minimizes, and transport weights make the batch subgradient exactly unbiased for  $\nabla F_p$  (Lemma 5).
2. *A coverage law and a severity measure* (Cor. 4). Feasibility is Hall-type and reduces on singletons to  $p_i \leq r_i$ , importance below observation rate, so the *infeasible mass*  $\nu \triangleq \sum_{i:p_i > r_i} p_i$  is the natural severity axis, computable before training.
3. *A two-term decomposition that locates the crossover* (Thm. 8, Sec. 5). Past the Hall condition the excess splits into an *alignment gain*  $\Delta_{\text{floor}}$ , the distance between the two optima each rule actually converges to, and an optimization residual. The gain is exact and instance-fixed; the residual is paid by the estimator. We sweep eleven skews and watch the two cross.
4. *Risk-averse layer* (Sec. 4). Mean-CVaR over groups composes with the transport weights without disturbing unbiasedness, at rate  $O(\kappa(\alpha, \gamma)/\sqrt{T})$ ,  $\kappa(\alpha, \gamma) = \gamma + (1 - \gamma)/\alpha$ .

**Related work.** Reduction methods for group fairness return randomized mixtures or carry an approximation constant tied to a fixed class [11, 12, 13]. Distributionally robust

and worst-group objectives [3, 2, 28] change the functional rather than the sampler and are complementary: our CVaR layer is one such functional, obtained on top of transport rather than instead of it. Importance sampling for SGD [16, 17] corrects variance, not the identity of the minimized objective.

## 2 PROBLEM FORMULATION

**Groups and objectives.** Let  $\mathcal{X} \subseteq \mathbb{R}^d$  be a feature space and  $\mathcal{Y}$  a label space, and let the population be partitioned into  $N$  groups indexed by  $i \in [N]$  (demographic categories, devices, environments) with *prevalences*  $\pi \in \Delta^{N-1}$ ,  $\pi_i = \mathbb{P}[\text{group } i]$ . Group  $i$  carries a distribution  $\mathcal{D}_i$  on  $\mathcal{X} \times \mathcal{Y}$  and the risk

$$f_i(\theta) \triangleq \mathbb{E}_{(X,Y) \sim \mathcal{D}_i} [\ell(m(X; \theta), Y)], \quad \theta \in \mathcal{C} \subseteq \mathbb{R}^d,$$

where  $m(\cdot; \theta)$  is the predictor and  $\ell$  the loss. Empirical risk minimization weights the groups by prevalence,  $F_\pi = \sum_i \pi_i f_i$ , and a majority group therefore dominates the trained model. The remedy is to fix an *importance distribution*  $p \in \Delta^{N-1}$  and minimize

$$F_p(\theta) \triangleq \sum_{i=1}^N p_i f_i(\theta), \quad \theta^* \in \arg \min_{\theta \in \mathcal{C}} F_p(\theta). \quad (1)$$

Uniform  $p$  is the category-uniform objective,  $p_i \propto 1/\pi_i$  the inverse-prevalence one, and  $p$  may equally encode a fairness constraint or be imposed externally [4, 5]. Throughout,

**Assumption 1.**  $\mathcal{C}$  is convex and compact with diameter  $D$ ; each  $f_i$  is convex on  $\mathcal{C}$  and  $0 \leq f_i \leq M$ ; and every stochastic subgradient  $g_i(\theta)$  of  $f_i$  returned by the sampler satisfies  $\mathbb{E}[\|g_i(\theta)\|^2] \leq G^2$  on  $\mathcal{C}$ .

These are the assumptions of [7, 6] restated for groups; they place no restriction on  $p$  or on  $\pi$ .

**Prevalence versus importance.** Minibatch SGD does not see  $p$ . At step  $t$  the sampler draws a batch  $B^t$  of  $b$  examples; write

$$A^t \triangleq \{i \in [N] : B^t \text{ contains an example from group } i\}$$

for the *active set* of groups, let  $\{A_j\}_{j \in [M]}$  enumerate its support, and let  $q_j \triangleq \mathbb{P}[A^t = A_j]$  be the *composition law*. The observation rate of group  $i$  is  $r_i \triangleq \mathbb{P}[i \in A^t] = \sum_{j: i \in A_j} q_j$ ; under i.i.d. sampling  $r_i = 1 - (1 - \pi_i)^b$ , so a group whose prevalence is below the resolution of the batch is absent from most batches whatever  $p$  says. This is federated learning under partial participation:  $A^t$  is the set of participating clients,  $q$  the availability distribution,  $p$  the importance distribution [6, 7]. We treat  $q$  as known; otherwise it is estimated offline by simulating the sampler, as in our experiments.

**What group-blind SGD minimizes.** The group-blind update averages the per-example subgradients of  $B^t$ , which assigns group  $i \in A^t$  the weight  $\hat{\pi}_i^t$ , its share of the batch, and assigns absent groups zero. Averaging first over the batch given  $A^t$  and then over  $A^t \sim q$  is the hierarchical experiment of [7, Eq. 4], and it names the function SGD is in fact minimizing.

**Proposition 2** (Surrogate objective). *Under Assumption 1, group-blind minibatch SGD is a stochastic subgradient method on  $F_{\tilde{p}}$  with*

$$\tilde{p}_i = \sum_{j: i \in A_j} q_j \mathbb{E}[\hat{\pi}_i^t | A^t = A_j], \quad (2)$$

a probability vector. For i.i.d. sampling  $\tilde{p} = \pi$ ; for uniform weighting within the batch  $\tilde{p}_i = \sum_{j: i \in A_j} q_j / |A_j|$ , the surrogate of agnostic FedAvg [6, Eq. 2]. Consequently, for every  $\theta \in \mathcal{C}$ ,

$$\|F_p(\theta) - F_{\tilde{p}}(\theta)\| \leq M \|p - \tilde{p}\|_1, \quad (3)$$

and any method that minimizes  $F_{\tilde{p}}$  exactly incurs the excess  $F_p(\theta_{\tilde{p}}^*) - F_p(\theta^*) \leq 2M \|p - \tilde{p}\|_1$  on (1), which does not vanish with  $T$ .

*Proof.* Conditionally on  $A^t = A_j$  the within-group averages are unbiased for  $\partial f_i$ , so  $\mathbb{E}[\text{batch subgradient} | A^t = A_j] = \sum_{i \in A_j} \mathbb{E}[\hat{\pi}_i^t | A_j] \partial f_i(\theta)$ ; averaging over  $q$  gives  $\sum_i \tilde{p}_i \partial f_i = \partial F_{\tilde{p}}$ , and  $\sum_i \tilde{p}_i = \sum_j q_j \sum_{i \in A_j} \mathbb{E}[\hat{\pi}_i^t | A_j] = 1$ . Both bounds follow from  $0 \leq f_i \leq M$  and the two-point inequality  $F_p(\theta_{\tilde{p}}^*) - F_p(\theta^*) \leq [F_p - F_{\tilde{p}}](\theta_{\tilde{p}}^*) + [F_{\tilde{p}} - F_p](\theta^*)$ .  $\square$

Proposition 2 says that the choice of  $p$  in (1) is decorative unless the aggregation step is changed: the iterates converge, at the usual rate, to the minimizer of a function whose weights are induced by the sampler rather than chosen, and the bias  $p - \tilde{p}$  survives every stepsize schedule. Rebalancing the sampler is only a partial fix:  $\tilde{p}_i \leq r_i$  always, so a group with  $p_i > r_i$  cannot be sampled into enough batches however the sampler is tilted, while upscaling the examples that do appear reintroduces an amplitude distortion that destabilizes the iterates [6]. The question is the one posed for federated aggregation in [6], read on batches:

*Can the per-batch group weights be chosen so that SGD driven by the composition law  $q$  converges for the objective defined by  $p$ ?*

The next section answers it, and characterizes exactly when the answer is yes.

## 3 PROPOSED APPROACH: ALIGNMENT BY MASKED OPTIMAL TRANSPORT

**Weights as a coupling.** Let  $Y[i, j] \geq 0$  be the weight assigned to group  $i$  when the active set is  $A_j$ , with  $Y[i, j] = 0$  for  $i \notin A_j$ , and let  $T[i, j] \triangleq q_j Y[i, j]$  be the joint mass carried from composition  $j$  to group  $i$ . Requiring that the weights be a probability vector on each realized composition, and that each group receive in expectation the mass  $p$  assigns it, is the *masked optimal transport* (MOT) feasibility problem of [6],

$$T \mathbf{1}_M = p, \quad \mathbf{1}_N^\top T = q, \quad T \geq 0, \quad \text{supp}(T) \subseteq E, \quad (4)$$

with mask  $E = \{(i, j) : i \in A_j\}$  encoding that a group absent from the batch cannot be weighted. Unlike classical transport [25, 20] there is no cost to minimize: any coupling respecting (4) serves, and (4) is the Kantorovich problem with the indicator cost  $C[i, j] = \infty$  off  $E$ . Fig. 1 is the picture.

**When alignment is possible.** Existence of a coupling is a supply-demand question on the bipartite graph  $E$ , and is settled by flow duality.




Figure 1: The mask  $E$ . Batch compositions carry the mass  $q$  they occur with; groups must receive the mass  $p$  the objective assigns them. Group 4 is reachable only through  $A_4$ , so exact alignment requires  $p_4 \leq q_4$ : the Hall condition of Thm. 3 read on the singleton  $I = \{4\}$ . Shrinking the batch deletes  $A_4$  from the support and makes  $p_4$  unreachable.

**Theorem 3** (Feasibility, [6, Thm. 1]). *For  $I \subseteq [N]$  let  $\mathcal{N}(I) = \{j : \exists i \in I, (i, j) \in E\}$  and  $\mathcal{S}(I) = \{j : A_j \subseteq I\}$ . Problem (4) is feasible if and only if*

$$\sum_{j \in \mathcal{S}(I)} q_j \leq \sum_{i \in I} p_i \leq \sum_{j \in \mathcal{N}(I)} q_j \quad \forall I \subseteq [N]. \quad (5)$$

Proofs are in the appendix. Read on singletons, (5) becomes a statement one can act on before training starts.

**Corollary 4** (Coverage law and severity). *Feasibility of (4) requires  $\mathbb{P}[A^t = \{i\}] \leq p_i \leq r_i$  for every  $i$ , with  $r_i = \mathbb{P}[i \in A^t]$  the observation rate of Sec. 2. Writing*

$$\nu \triangleq \sum_{i : p_i > r_i} p_i \quad (6)$$

for the infeasible mass,  $\nu = 0$  iff the upper constraints hold. Under i.i.d. sampling of  $b$  examples the upper constraint reads  $b \geq \log(1 - p_i) / \log(1 - \pi_i)$ ; under  $K$ -of- $N$  uniform group sampling it reads  $p_i \leq K/N$ .

**Severity.** Importance may not exceed observation rate: a rare group is reachable only if the batch is large enough to contain it often enough. The infeasible mass  $\nu$  is the axis we sweep: at  $\nu = 0$  exact alignment is available, and as  $\nu$  grows more of  $p$  sits outside the set of row marginals the mask can reach. A *feasible* and an *infeasible* regime are two points on this continuum, and Sec. 5 sweeps it instead.

**Computing the plan.** Given feasibility, a coupling is obtained by iterative proportional fitting (Sinkhorn scaling) on the masked support [18, 19, 20]: starting from any  $T^{(0)} > 0$  on  $E$  with  $\mathbf{1}^\top T^{(0)} = q$ , alternate

$$T \leftarrow \text{diag}(p \oslash T \mathbf{1}) T, \quad T \leftarrow T \text{diag}(q \oslash T^\top \mathbf{1}) \quad (7)$$

until  $\|T \mathbf{1} - p\|_1 \leq \varepsilon$ . Each sweep costs  $O(|E|)$  with  $|E| \leq NM$  and  $M$  the number of *realized* compositions rather than  $2^N$ ; the plan is computed once, offline, and the training loop adds one lookup and one multiply per group per step. Under (5) the iteration converges to a solution of (4) [18, 22], being block-coordinate ascent on the dual of the entropic problem  $\min\{\sum_E T \log T : T \mathbf{1} = p, \mathbf{1}^\top T = q\}$ . Since  $\mathbf{1}_N^\top T = q$  and  $Y = T \text{diag}(q)^{-1}$ , every column  $Y[:, j]$  is a probability vector on  $A_j$ , the property each bound below uses. Algorithm 1 is the resulting method;  $\gamma = 1$  is the risk-neutral rule analyzed here and  $\gamma < 1$  the risk-averse layer of Sec. 4.

---

**Algorithm 1** GAVOT: group-fair SGD by masked transport

---

**Require:** target  $p$ , composition law  $(q, \{A_j\}, E)$ , step  $\eta$ , horizon  $T$ , CVaR parameters  $(\alpha, \gamma)$  with  $\gamma = 1$  for the risk-neutral method

```

1:  $(T, Y) \leftarrow$  Sinkhorn scaling for  $(p, q, E)$   $\triangleright$  offline, once
2:  $\theta^0 \in \mathcal{C}$ ,  $t^0 \in \mathbb{R}$ 
3: for  $t = 1, \dots, T$  do
4:   draw batch  $B^t$ ; set  $A^t$ , let  $j(t)$  be s.t.  $A_{j(t)} = A^t$ 
5:    $g_i^t \leftarrow$  stochastic subgradient of  $f_i$  from  $B^t \cap$  group  $i$ 
6:    $w_i^t \leftarrow Y[i, j(t)](\gamma + \frac{1-\gamma}{\alpha} \mathbf{1}\{\hat{f}_i^t > t^{t-1}\})$ 
7:    $\theta^t \leftarrow \Pi_{\mathcal{C}}(\theta^{t-1} - \eta \sum_{i \in A^t} w_i^t g_i^t)$ 
8:    $t^t \leftarrow t^{t-1} - \eta(1 - \gamma)(1 - \frac{1}{\alpha} \sum_{i \in A^t} Y[i, j(t)] \mathbf{1}\{\hat{f}_i^t > t^{t-1}\})$ 
9: end for
Ensure:  $\bar{\theta}_T = \frac{1}{T} \sum_{t=1}^T \theta^t$ 

```

---

**Analysis.** Everything rests on one identity, [6, Lem. 1] moved from model averaging to gradient averaging. Let  $g_i(\theta)$  be the stochastic subgradient of  $f_i$  computed from  $B^t \cap$  group  $i$  and  $\hat{g}(\theta) = \sum_{i \in A^t} Y[i, j(t)] g_i(\theta)$  the transport-weighted batch subgradient of Algorithm 1. Transport removes bias; it is not a variance reduction device.

**Lemma 5** (Unbiasedness). *Let  $T$  satisfy (4). Then  $\mathbb{E}[\hat{g}(\theta) | \theta] \in \partial F_p(\theta)$  for every  $\theta \in \mathcal{C}$ .*

*Proof.* Conditionally on  $A^t = A_j$ ,  $\mathbb{E}[g_i(\theta)] \in \partial f_i(\theta)$  for  $i \in A_j$ . Averaging over  $A^t \sim q$ ,  $\mathbb{E}[\hat{g}] = \sum_j q_j \sum_{i \in A_j} Y[i, j] \partial f_i(\theta) = \sum_i (\sum_j T[i, j]) \partial f_i(\theta) = \sum_i p_i \partial f_i(\theta)$ , the row-marginal constraint being used in the last step.  $\square$

**Lemma 6** (Second moment, free of coverage).  $\|\hat{g}(\theta)\|^2 \leq G^2$  for every  $\theta \in \mathcal{C}$ , with no dependence on  $|A^t|$  or on  $N$ .

*Proof.*  $Y[:, j]$  is a probability vector on  $A_j$ , so  $\|\sum_i Y[i, j] g_i\|^2 \leq \sum_i Y[i, j] \|g_i\|^2$  by convexity of  $\|\cdot\|^2$ ; take expectations and use  $\sum_j T[i, j] = p_i$ :  $\mathbb{E}\|\hat{g}\|^2 \leq \sum_i p_i \mathbb{E}\|g_i\|^2 \leq G^2$ .  $\square$

**Theorem 7** (Rate). *Under Assumption 1, if (4) is feasible then Algorithm 1 with  $\gamma = 1$  and  $\eta = D/(G\sqrt{T})$  satisfies*

$$\mathbb{E}[F_p(\bar{\theta}_T)] - F_p(\theta^*) \leq \frac{DG}{\sqrt{T}}. \quad (8)$$

The constant in (8) is that of full-participation SGD and is independent of how many groups appear in a batch: a composition covering two groups yields the same bound as one covering all  $N$ , though with larger realized variance. The proof uses only the two lemmas and never refers to  $|A^t|$ , which is why the rate of [6] transfers verbatim from aggregation of models to aggregation of gradients.

**Outside the feasible region.** When (5) fails, (4) has no solution and the iteration (7) does not diverge: its iterates accumulate at the coupling whose row marginal  $\hat{p} \triangleq T \mathbf{1}$  is the KL-closest to  $p$  over the reachable polytope  $\mathcal{R}(q, E) = \{T \mathbf{1} : T \geq 0, \mathbf{1}^\top T = q, \text{supp}(T) \subseteq E\}$  [22, 21, 20]. The column marginal is met exactly, so Lemmas 5 and 6 hold with  $p$  replaced by  $\hat{p}$ : Algorithm 1 is then unbiased SGD on the surrogate  $F_{\hat{p}}$ , the closest objective the sampler can

reach, and the loss on (1) is a bias of exactly the shape of Proposition 2. On singletons  $\hat{p}_i \leq r_i$ , so the groups with  $p_i > r_i$  are truncated at their observation rate and their missing mass,  $\nu$  up to renormalization, is redistributed.

**Theorem 8** (Infeasible regime). *Let  $\hat{p}$  be the row marginal of the Sinkhorn limit,  $\theta^\dagger \in \arg \min_C F_{\hat{p}}$ , and write  $\Phi(w) \triangleq F_p(\arg \min_C F_w)$  for the  $p$ -risk of the minimizer of the  $w$ -weighted risk. Then*

$$\mathbb{E}[F_p(\bar{\theta}_T)] - F_p(\theta^\dagger) = \underbrace{\Phi(\hat{p}) - \Phi(p)}_{\text{alignment floor}} + \underbrace{\varepsilon_T(\hat{p})}_{\text{optimization residual}}, \quad (9)$$

where  $\varepsilon_T(\hat{p}) \triangleq \mathbb{E}F_p(\bar{\theta}_T) - F_p(\theta^\dagger)$  satisfies  $\mathbb{E}F_{\hat{p}}(\bar{\theta}_T) - F_{\hat{p}}(\theta^\dagger) \leq DG/\sqrt{T}$  and  $\varepsilon_T(\hat{p}) \rightarrow 0$  whenever  $\bar{\theta}_T \rightarrow \theta^\dagger$ . Group-blind averaging obeys (9) with  $\tilde{p}$  in place of  $\hat{p}$ . Writing  $\Delta_{\text{floor}} \triangleq \Phi(\tilde{p}) - \Phi(\hat{p})$  for the alignment gain, transport improves on group-blind averaging exactly when  $\Delta_{\text{floor}} > \varepsilon_T(\hat{p}) - \varepsilon_T(\tilde{p})$ , and every floor is bounded by  $\Phi(w) - \Phi(p) \leq 2M\|p - w\|_1$ .

Theorem 8 replaces the usual bias bound [6] by an exact decomposition, both terms of which are available before training:  $\hat{p}$  and  $\tilde{p}$  from  $(p, q, E)$ , and  $\Phi$  exactly on instances with a closed-form or convex frontier. The gain is fixed by the instance and decides the sign of the comparison; the residual is paid by the estimator and, as Sec. 5 shows, is a stepsize artifact. The  $\ell_1$  bound is the weaker object and we report both, because it turns out too loose to discriminate.

#### 4 A RISK-AVERSE LAYER OVER GROUPS

Exact alignment delivers the  $p$ -weighted risk, an average, which can hide a group. Where the requirement is on the worst group instead, we compose a mean-CVaR functional over groups on top of the transport weights [7, 26]:

$$F_p^{\alpha, \gamma}(\theta) \triangleq \gamma F_p(\theta) + (1 - \gamma) \text{CVaR}_\alpha^p(f(\theta)), \quad \alpha, \gamma \in (0, 1], \quad (10)$$

where  $\text{CVaR}_\alpha^p$  is the conditional value at risk of the group-loss vector under  $p$ . The Rockafellar–Uryasev representation turns (10) into a joint minimization over  $(\theta, t)$ ,

$$\Phi(\theta, t) = \gamma \sum_i p_i f_i(\theta) + (1 - \gamma) \left( t + \frac{1}{\alpha} \sum_i p_i [f_i(\theta) - t]_+ \right), \quad (11)$$

jointly convex, with subgradients in both arguments again  $p$ -weighted sums of per-group quantities, so Lemma 5 applies unchanged. This is why the layers compose rather than interfere: transport fixes *which* distribution the expectation is under, CVaR changes *what* is averaged. Lines 6 and 8 of Algorithm 1 implement (11).

**Theorem 9** (Rate under mean-CVaR). *Under the hypotheses of Theorem 7, Algorithm 1 with  $\eta = \Theta(1/\sqrt{T})$  satisfies  $\mathbb{E}[F_p^{\alpha, \gamma}(\bar{\theta}_T)] - \min F_p^{\alpha, \gamma} = O(\kappa(\alpha, \gamma)/\sqrt{T})$  with  $\kappa(\alpha, \gamma) = \gamma + (1 - \gamma)/\alpha$ . In the infeasible regime the bound of Theorem 8 acquires the same factor on its statistical term, the bias term being unchanged.*

Both  $\gamma = 1$  and  $\alpha = 1$  recover Theorem 7 exactly, so the grid contains the risk-neutral method along two edges, visible as the constant row and column of every heatmap below. Sending  $\alpha \downarrow 0$  drives (10) to the worst group at the cost  $\kappa(\alpha, \gamma) \rightarrow \infty$ : the parameter prices robustness in the convergence constant, not in the limit.

Table 1: Headline cells, tail-500 over 5 seeds, decaying stepsize. *full* sees every group every step. The fixed multiplier  $m/K$  diverges whenever observation rates are non-uniform.

| Instance                        | GAVOT         | unif.  | avg $m/K$ | full   |
|---------------------------------|---------------|--------|-----------|--------|
| Adult, prevalence ( $\nu=.60$ ) | <b>0.2269</b> | 0.2479 | 0.670     | 0.2073 |
| Adult, aligned ( $\nu=0$ )      | 0.2075        | 0.2077 | div.      | 0.2073 |
| IMDb, skewed ( $\nu=.31$ )      | <b>86.06</b>  | 91.63  | 8504      | 82.95  |
| IMDb, aligned ( $\nu=0$ )       | <b>84.28</b>  | 86.37  | div.      | 82.95  |




Figure 2: IMDb-Wiki severity sweep. Solid, measured tail-500 ( $\pm 1$  std); dashed, the exact floor  $\Phi(\cdot)$  each rule converges to. Left, constant stepsize: the optimization residual swamps the alignment gain and transport loses everywhere, even where the floors are seven units apart. Right, decaying stepsize: the residual collapses and the measured curves track their floors, so transport wins until the floors close at high  $\nu$ . The sign is set by which of the two terms of (9) dominates, not by the regime label.

**Remark 10** (Relation to constrained learning).  $\max_i f_i$  is the exact-penalty limit of  $\min f_0$  s.t.  $f_i \leq c_i$ , a fixed price on the largest violation matching penalized and constrained values once it exceeds the least  $\ell_1$  norm on the dual-optimal set [15]. The  $\alpha \downarrow 0$  edge of (10) is thus the scalarization a constrained specification selects, with  $\kappa$  its price in the rate.

#### 5 EXPERIMENTS

**Instances.** *Adult* income classification [33] (logistic regression,  $\eta=0.1$ ) with a uniform importance target over the five race groups, 100 race-homogeneous groups of 30 samples split 85/10/3/1/1 by prevalence, and *IMDb-Wiki* age regression [34] (linear regression on ResNet embeddings,  $\eta=10^{-2}$ ) over 100 identities with a cubically skewed importance target. Both use  $K = 3$  observed groups per step,  $H=5$  local steps per observed group before the weighted average, 4000 steps, 5 seeds, tail-500 statistics. Availability rates  $r$  are swept through a family indexed by  $\beta$ , from prevalence-proportional to importance-aligned on Adult and through increasing skew on IMDb-Wiki, giving eleven instances whose infeasible mass  $\nu$  of (6) ranges from 0 to 88%. On both the perturbation frontier is closed-form, so the floors  $\Phi(p)$ ,  $\Phi(\hat{p})$ ,  $\Phi(\tilde{p})$  of Theorem 8 are exact rather than estimated. Besides group-blind averaging we include the fixed multiplier  $m/K$ , the rescaling that assumes a uniform observation rate.

**E1: the headline cells.** Table 1 reports the four instances a practitioner would recognize. Transport wins on three: Adult under prevalence-proportional availability, 0.2269 against 0.2479 ( $t=45$ ), and both IMDb cells, 86.06 against 91.63 ( $t=21$ ) and 84.28 against 86.37 ( $t=13$ ). On Adult with importance-aligned availability the two coincide at the oracle value, which is what Proposition 2 predicts when  $\tilde{p} = p$ : there is nothing to correct and transport does not manufacture a difference. The fixed multiplier diverges on both aligned instances and is far off on the other two, so

Table 2: IMDb-Wiki severity sweep, decaying stepsize.  $\Delta_{\text{floor}}$  is the alignment gain of Theorem 8, *gain* the measured advantage of transport,  $t$  its Welch statistic (population std). The excess residual transport pays averages 1.9, and the sign flips exactly where  $\Delta_{\text{floor}}$  crosses it. The five Adult skews all sit above their crossover and are omitted.

| $\beta$ | $\nu$ | $\Delta_{\text{floor}}$ | gain         | $t$  |
|---------|-------|-------------------------|--------------|------|
| 0       | .31   | 6.99                    | <b>+5.57</b> | 20.7 |
| 0.5     | .59   | 5.71                    | <b>+3.46</b> | 11.3 |
| 1.0     | .70   | 3.67                    | <b>+2.05</b> | 5.4  |
| 1.5     | .77   | 2.56                    | +0.58        | 1.1  |
| 2.0     | .82   | 1.62                    | −0.03        | −0.1 |
| 3.0     | .88   | 0.02                    | −2.28        | −7.9 |

naive rescaling is no alternative to solving (4).

**E2: the crossover, and where the sign comes from.** Table 2 and Fig. 2 are the paper’s main evidence. On Adult, whose five skews are omitted from the table, the measured gain tracks  $\Delta_{\text{floor}}$  within 4% at three of five (.0134 against .0129, .0050 against .0049, .0002 against .0003) and exceeds it at the two most infeasible (.0210 against .0110, .0181 against .0144), group-blind averaging sitting further above its floor. On IMDb-Wiki the gain is uniformly below the floor by an excess optimization residual that is close to constant in  $\beta$ , mean 1.9 in MSE units, and the measured sign flips precisely where  $\Delta_{\text{floor}}$  crosses that value, between  $\beta=1.5$  (2.56, gain +0.58) and  $\beta=2.0$  (1.62, gain −0.03). This is Theorem 8 read at finite  $T$ : the sign of the comparison is a quantity one computes in advance.

**E3: the  $\ell_1$  bound is too loose to be the criterion.** The distances  $\|p - \hat{p}\|_1$  and  $\|p - \tilde{p}\|_1$  are also computed on every instance, and  $\|p - \hat{p}\|_1 < \|p - \tilde{p}\|_1$  holds in all eleven, including the three where transport loses or ties. The  $\ell_1$  form of Theorem 8 therefore never predicts a loss and cannot serve as the criterion; the exact floors can.

**E4: the residual is a stepsize artifact.** The left panel of Fig. 2 runs the identical instances at constant stepsize. There transport loses on all six IMDb skews, by as much as 7.9, even at  $\nu=.31$  where the floors are 6.99 apart. The transport weights  $Y[\cdot, j]$  concentrate mass on rarely observed groups, so although Lemma 6 bounds the second moment identically for both rules, the variance about the mean is larger, and at constant stepsize the iterates sit on a noise floor proportional to it. Under  $\eta_t = \eta/(1 + t/1000)$  that floor decays, the measured curves fall onto their dashed floors, and every sign reversal attributable to the stepsize disappears. The residual is a property of the schedule, not of the estimator.

**E5: risk aversion, and where it does not help.** Sweeping  $(\alpha, \gamma)$  on a  $10 \times 10$  grid, the overall objective is minimized on the  $\gamma = 1$  edge on every instance, which is  $\kappa(1, \gamma) = 1$  made visible and confirms that risk aversion costs the average. The one place the CVaR layer earns its keep is the worst race on Adult under prevalence availability, where GAVOT+CVaR at (0.2, 0.7) reaches 0.3624 against 0.3679 for GAVOT and 0.3994 for group-blind averaging. Per race, transport improves four of five (Black .179 vs .200, Asian-Pac .368 vs .399, Am-Ind .175 vs .189, Other .073 vs .119) and reweights only the majority group (.339 vs .332), which is what reweighting toward a category-uniform target is supposed to

do.

## 6 CONCLUSION

Group-fair training fails in a specific and fixable place: the aggregation step, blind to the weighting the objective declares. Reading a minibatch as a partial-participation event makes the fix the masked transport problem that corrects availability skew in federated learning, with the same feasibility characterization, an  $O(1/\sqrt{T})$  guarantee free of batch coverage, and a coverage law whose infeasible mass  $\nu$  orders instances by severity. Past feasibility the excess splits cleanly, into an alignment gain fixed by the instance and an optimization residual paid by the estimator, and a sweep over eleven skews shows the two crossing where the decomposition says they will. Transport is worth applying while the gain exceeds the residual, both are computable before training, and the residual is a stepsize artifact a decaying schedule removes. The nonconvex case remains open.

## REFERENCES

- [1] J. Buolamwini and T. Gebru, “Gender shades: Intersectional accuracy disparities in commercial gender classification,” in *Proc. Conf. Fairness, Accountability and Transparency (FAT\*)*, 2018, pp. 77–91.
- [2] T. Hashimoto, M. Srivastava, H. Namkoong, and P. Liang, “Fairness without demographics in repeated loss minimization,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2018, pp. 1929–1938.
- [3] S. Sagawa, P. W. Koh, T. Hashimoto, and P. Liang, “Distributionally robust neural networks for group shifts,” in *Proc. Int. Conf. Learning Representations (ICLR)*, 2020.
- [4] M. Mohri, G. Sivek, and A. T. Suresh, “Agnostic federated learning,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2019, pp. 4615–4625.
- [5] T. Li, M. Sanjabi, A. Beirami, and V. Smith, “Fair resource allocation in federated learning,” in *Proc. Int. Conf. Learning Representations (ICLR)*, 2020.
- [6] H. Rahimi and D. Kalogerias, “FedAVOT: Exact distribution alignment in federated learning via masked optimal transport,” in *Proc. IEEE Int. Conf. Acoustics, Speech and Signal Processing (ICASSP)*, 2026.
- [7] H. Rahimi and D. Kalogerias, “Convergence of agnostic federated averaging,” in *Proc. IEEE Int. Workshop Computational Advances in Multi-Sensor Adaptive Processing (CAMSAP)*, 2025.
- [8] J. Wang, Q. Liu, H. Liang, G. Joshi, and H. V. Poor, “Tackling the objective inconsistency problem in heterogeneous federated optimization,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 7611–7623.
- [9] Y. J. Cho, J. Wang, and G. Joshi, “Towards understanding biased client selection in federated learning,” in *Proc. Int. Conf. Artificial Intelligence and Statistics (AISTATS)*, 2022, pp. 10351–10375.
- [10] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Proc. Int. Conf. Artificial Intelligence and Statistics (AISTATS)*, 2017, pp. 1273–1282.
- [11] A. Agarwal, A. Beygelzimer, M. Dudík, J. Langford, and H. Wallach, “A reductions approach to fair classification,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2018, pp. 60–69.
- [12] A. Cotter, H. Jiang, and K. Sridharan, “Two-player games for efficient non-convex constrained optimization,” in *Proc. Int. Conf. Algorithmic Learning Theory (ALT)*, 2019, pp. 300–332.
- [13] L. F. O. Chamon and A. Ribeiro, “Probably approximately correct constrained learning,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 16722–16735.
- [14] L. F. O. Chamon, S. Paternain, M. Calvo-Fullana, and A. Ribeiro, “Constrained learning with non-convex losses,” *IEEE Trans. Inf. Theory*, vol. 69, no. 3, pp. 1739–1760, 2023.
- [15] H. Rahimi, S. Pougkakiotis, and D. Kalogerias, “An exact penalty method for constrained statistical learning over universal RKHSs,” 2026, submitted.
- [16] M. Rizk, S. Vlaski, and A. H. Sayed, “Optimal importance sampling for federated learning,” in *Proc. IEEE Int. Conf. Acoustics, Speech and Signal Processing (ICASSP)*, 2021, pp. 3110–3114.
- [17] A. Katharopoulos and F. Fleuret, “Not all samples are created equal: Deep learning with importance sampling,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2018, pp. 2525–2534.
- [18] R. Sinkhorn, “Diagonal equivalence to matrices with prescribed row and column sums,” *Amer. Math. Monthly*, vol. 74, no. 4, pp. 402–405, 1967.
- [19] M. Cuturi, “Sinkhorn distances: Lightspeed computation of optimal transport,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 26, 2013, pp. 2292–2300.
- [20] G. Peyré and M. Cuturi, *Computational Optimal Transport*. Now Publishers, 2019.
- [21] C. Gietl and F. P. Reffel, “Accumulation points of the iterative proportional fitting procedure,” *Statistics*, vol. 47, no. 6, pp. 1321–1335, 2013.
- [22] I. Csiszár, “ $I$ -divergence geometry of probability distributions and minimization problems,” *Ann. Probab.*, vol. 3, no. 1, pp. 146–158, 1975.
- [23] P. Hall, “On representatives of subsets,” *J. London Math. Soc.*, vol. 10, no. 1, pp. 26–30, 1935.
- [24] L. R. Ford and D. R. Fulkerson, “Maximal flow through a network,” *Canad. J. Math.*, vol. 8, pp. 399–404, 1956.
- [25] C. Villani, *Optimal Transport: Old and New*. Springer, 2009.
- [26] R. T. Rockafellar and S. Uryasev, “Optimization of conditional value-at-risk,” *J. Risk*, vol. 2, no. 3, pp. 21–41, 2000.
- [27] S. Curi, K. Y. Levy, S. Jegelka, and A. Krause, “Adaptive sampling for stochastic risk-averse learning,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 1036–1047.
- [28] D. Levy, Y. Carmon, J. C. Duchi, and A. Sidford, “Large-scale methods for distributionally robust optimization,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 8847–8860.
- [29] H. Namkoong and J. C. Duchi, “Variance-based regularization with convex objectives,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 30, 2017, pp. 2971–2980.
- [30] A. Nemirovski, A. Juditsky, G. Lan, and A. Shapiro, “Robust stochastic approximation approach to stochastic programming,” *SIAM J. Optim.*, vol. 19, no. 4, pp. 1574–1609, 2009.
- [31] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith, “Federated optimization in heterogeneous networks,” in *Proc. Machine Learning and Systems (MLSys)*, 2020, pp. 429–450.
- [32] S. P. Karimireddy, S. Kale, M. Mohri, S. J. Reddi, S. U. Stich, and A. T. Suresh, “SCAFFOLD: Stochastic controlled averaging for federated learning,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2020, pp. 5132–5143.
- [33] B. Becker and R. Kohavi, “Adult,” UCI Machine Learning Repository, 1996.
- [34] R. Rothe, R. Timofte, and L. Van Gool, “Deep expectation of real and apparent age from a single image without facial landmarks,” *Int. J. Comput. Vis.*, vol. 126, no. 2–4, pp. 144–157, 2018.
- [35] P. Theodoropoulos, K. E. Nikolakakis, and D. Kalogerias, “Federated learning under restricted user availability,” in *Proc. IEEE Int. Conf. Acoustics, Speech and Signal Processing (ICASSP)*, 2024.